

# On Extending a Full-Sharing Design with Batched Scheduling

Miguel Areias and Ricardo Rocha

CRACS & INESC TEC and Faculty of Sciences, University of Porto  
Rua do Campo Alegre, 1021, 4169-007 Porto, Portugal  
{miguel-areias,ricroc}@dcc.fc.up.pt

**Abstract.** **Keywords:** Multithreading, Tabling, Concurrency, Batched Scheduling

## 1 Introduction

## 2 Background

### 2.1 No-Sharing, Subgoal-Sharing and Full-Sharing Designs

### 2.2 Batched Scheduling

The decision about the evaluation flow is determined by the *scheduling strategy*. Different strategies may have a significant impact on performance, and may lead to a different ordering of solutions to the query goal. Arguably, the two most successful tabling scheduling strategies are *local scheduling* and *batched scheduling* [3].

Local scheduling strategy schedules the evaluation of a program in a breath-first manner. It favors the backtracking first with completion instead of the forward execution, leaving the consumption of answers for last. Thus, it only allows a Cluster of Dependent Subgoals (CDS) to return answers only after the completion point has been reached [3]. In other words, the local scheduling tries to keep a CDS as minimal as possible. When new answers are found, they are added to the table space and the computation fails as consequence, tabled subgoals inside a CDS propagate their answers to outside the CDS only after its completion point is found. Local scheduling causes a sooner completion of subgoals, which creates less complex dependencies between them.

On the other hand, batched scheduling schedules the evaluation of a program in a depth-first manner. It favors the forward execution first instead of backtracking, leaving the consumption of answers and completion for last. It thus tries to delay the need to move around the search tree by batching the return of answers. When new answers are found for a particular tabled subgoal, they are added to the table space and the execution continues. For some situations, this results in creating dependencies to older subgoals, therefore enlarging the current CDS [4] and delaying the completion point to an older generator node.

### 3 Full-Sharing with Batched Scheduling


#### 3.1 Our Approach

The key idea of our approach, which we named Privately-consumed Answer Chaining (PAC), is then to extend the *FS* design with batched scheduling, by chaining privately for each subgoal call, the answers that were already consumed by a thread. Since the procedure is private, it will only affect the thread that is doing it. At the end, when the evaluation is complete, i.e., when a subgoal call is marked as complete, we put one of the private chain as public, so that from that point on all threads can use that chain in complete (only reading) mode.

Figure 1 shows the key data structures for supporting the implementation of the PAC procedure during the evaluation of a tabled subgoal call  $P_{i,j}$  using the *FS* design. The *FS* design, uses a subgoal entry data structure to store common information for a subgoal call and a subgoal frame (*SF*) data structure to store private information about the execution of each thread. The PAC procedure works at the subgoal frame level, which is private to each thread.

Figure 1(a) shows then a situation where two threads,  $T_1$  and  $T_{k-1}$ , are sharing the same subgoal entry call  $P_{i,j}$  when the subgoal is still under evaluation, i.e., the subgoal is not yet complete. The current state of the evaluation shows an answer trie with 3 answers found for the subgoal call  $P_{i,j}$ . For the sake of simplicity, we are omitting the internal answer trie nodes and we are only showing the leaf nodes  $LN_1$ ,  $LN_2$  and  $LN_3$ , in the figure. With the PAC procedure, the leaf nodes are not chained in the *AT* data structure. Now, the chaining process is done privately, and for that, we use the subgoal frame structure of each thread. On the subgoal frame structure we added a new field, called *consumer*, to store the answers found within the execution of the thread. In order to minimize the impact of the *PAC* optimization, each node within the new consumer answer structure has two fields: (i) an entry pointer, which points to the corresponding leaf node in the answer trie data structure; (ii) a next pointer to chain the answers within the consumer structure. To maintain a good performance, when the number of nodes exceeds a certain threshold, we use a hash trie mechanism design similar to the one presented in the work [2]. However, since this mechanism is private to each thread, it does not require any of the tools that were necessary to support concurrency. In particular, on each hash trie level, we have removed the tools necessary to support concurrency, such as useless pointers and compare-and-swap operations. We have chosen this hashing mechanism, because it showed a good balance between lookup and insert operations [2], but the major reason was mostly because of the integration in the TabMalloc memory allocator [1].

Going back to Figure 1(a), the consumer answer structures represent then two different situations where threads can be evaluating a subgoal call. Thread  $T_1$  has only found one answer and it is using a direct consumer answer chaining to access the node  $LN_1$ . Thread  $T_{k-1}$  was already found three answers for the subgoal call and it is already using the hash trie mechanism within its consumer answer structure. The consumer nodes are chained between themselves, thus



**Fig. 1.** The *FS* design with the PAC procedure - (a) private chaining and (b) public chaining

that consumer nodes belonging to thread  $T_{k-1}$  can consume the answers as in the original mechanism.

Figure 1(b) shows the state of the subgoal call after completion (recall that after completion of a subgoal call, the threads use loader nodes to consume the answers). When a thread  $T$  completes a subgoal call, it frees its private consumer structures, but before doing that, it checks whether another thread has already marked the subgoal as completed. If no other thread has done that, then thread  $T$  not only follows its private chaining mechanism as it would for freeing its private nodes, but also, follows the pointers to the answer trie leaf nodes in order to reproduce the chain inside the answer trie. Since this procedure is done inside a critical region, no more than one thread can be doing this chaining

process. Thus, in Figure 1(b), we are showing a situation where the subgoal call is completed and both threads  $T_1$  and  $T_{k-1}$  have already removed their consumer answer structures and chained the leaf nodes inside the answer trie.

### 3.2 Implementations details

At the implementation level, the major difference between local and batched scheduling is in the tabling operation *tabled new answer*, where we decide what to do when an answer is found during the evaluation. This operation checks whether a newly found answer is already in the corresponding answer trie structure and, if not, inserts it. For the NS and SS designs the support for batched scheduling is immediate, since the answer trie data structure is not shared among threads. The usage of batched scheduling with the FS design requires further support since with batched scheduling, answers are immediately propagated and we have to ensure that the propagation of an answer occurs on all subgoal calls one and only once. To do so, we take advantage of the private chaining procedure, presented in the previous subsection, as a way to keep, for every subgoal call of every thread, track of all the answers that were already propagated. This requires minor changes to the *tabled new answer* tabling operation. Algorithm 1 shows how we have extended the tabled new answer operation to support the FS design with batched scheduling.

---

**Algorithm 1** *tabled\_new\_answer*(answer *ANS*, subgoal frame *SF*)

---

```

1: leaf  $\leftarrow$  check_insert_answer_trie(ANS, SF)
2: if NS_design or SS_design then
3:   ... {without changes}
4: else {FS design}
5:   chain  $\leftarrow$  check_insert_consumer_chain(leaf, SF)
6:   if is_answer_marked_as_found(chain) = True then
7:     return failure
8:   else {the answer is new}
9:     mark_answer_as_found(chain)
10:    if local_scheduling_mode(SF) then
11:      return failure
12:    else {batched scheduling mode}
13:      return proceed

```

---

The algorithm receives two arguments: the new answer found during the evaluation (*ANS*) and the subgoal frame which corresponds to the call at hand (*SF*). The *NS\_design*, *SS\_design* and *FS\_design* macros define which table design is enabled.

The algorithm begins by checking/inserting the given *ANS* into the answer trie structure, which will return the leaf node for the path representing *ANS* (line 1). In line 2, it then tests whether one of the NS or SS designs are active, and in such a case, the algorithm is remains unchanged.

Otherwise, for the FS design (lines 4 to 13), it checks/inserts the given *leaf* node into the private consumer chain for the current thread, which will return the corresponding chain node. In line 6, it then tests whether the chain node already existed in the consumer chain, i.e., if it was inserted or not by the current check/insert operation in order to return failure (line 7), or it proceed with marking the answer *ANS* has found (line 9). At the end (lines 10 to 13), it returns failure if local scheduling is active (line 11), otherwise, the batched scheduling is active, thus it propagates the answer *ANS* (line 13).

## 4 Performance Analysis

We now present experimental results about the usage of the batched scheduling on the *NS*, *SS* and *FS* designs. For the sake of simplicity, for the *SS* and *FS* designs, we will be presenting only the results for the lock free *LF2* proposal ( $SS_{LF_2}$  and  $FS_{LF_2}$ ), since they were the ones that presented the lowest overheads in the previous chapters. For the  $FS_{LF_2}$  design, we will use it with the *PCC* procedure enabled.

Concerning the benchmarks, we will be using the same five sets of benchmarks presented before with the same number of runs per benchmark, the same formula to calculate the overhead ratios, and the same worst case scenario approach, where all threads begin with the same query goal. To put the results in perspective, we experimented with 1, 8, 16, 24 and 32 threads (the maximum number of cores available in our machine) with batched and local scheduling.

Table 1 shows the overhead ratios, when compared with the *NS* design with 1 thread (running with local scheduling, PtMalloc and without TabMalloc), for the *NS*,  $SS_{LF_2}$  and  $FS_{LF_2+PCC}$  designs (all running with TabMalloc and TcMalloc), when running 1, 8, 16, 24 and 32 threads with local and batched scheduling on the five sets of benchmarks (the results by set of benchmark can be seen in the Appendix. ??). For each design, the table has then two columns, a column with **Local** that shows results already presented in previous chapters for the local scheduling and a column with **Batched** that shows the new results with the batched scheduling. The overhead results presented in both **Local** and **Batched** columns use as base time the execution times presented in the *NS* column of the Table ??.

By observing Table 1, we can see that, for one thread, on average, local scheduling is slightly better than batched on the three designs. For the *NS* design we have 0.78 and 0.82, for the  $SS_{LF_2}$  design we have 0.84 and 0.90 and for the  $FS_{LF_2+PCC}$  design we have 1.30 and 1.46 average overhead ratios, for the local and batched scheduling strategies, respectively.

As we scale the number of threads, one can observe that, for the *NS* and  $SS_{LF_2}$  designs both scheduling strategies have similar minimum, average and maximum overhead ratios. For the  $FS_{LF_2+PCC}$  design, the best minimum overhead ratio is always for batched scheduling. The reader can observe on Appendix ?? that the minimum overhead values for 8, 16, 24 and 32 threads are given by the benchmark belonging to the model checking set (see Table ?? for

**Table 1.** Overhead ratios, when compared with the NS design with 1 thread (running with local scheduling, PtMalloc and without TabMalloc) for the NS,  $SS_{LF_2}$ ,  $FS_{LF_2+PCC}$  designs (with TabMalloc and TcMalloc), when running 1, 8, 16, 24 and 32 threads with local and batched scheduling on the five sets of benchmarks (best ratios by row and by design for the Minimum, Average and Maximum are in bold)

Threads	NS		SS <sub>LF<sub>2</sub></sub>		FS <sub>LF<sub>2</sub>+PCC</sub>		
	Local Batched		Local Batched		Local Batched		
1	Min	<b>0.53</b>	0.55	<b>0.54</b>	0.55	1.01	<b>0.95</b>
	Avg	<b>0.78</b>	0.82	<b>0.84</b>	0.90	<b>1.30</b>	1.46
	Max	1.06	<b>1.05</b>	<b>1.04</b>	<b>1.04</b>	<b>1.76</b>	2.33
	StD	0.15	0.14	0.17	0.16	0.22	0.44
8	Min	0.66	<b>0.63</b>	0.66	<b>0.63</b>	1.16	<b>0.99</b>
	Avg	<b>0.85</b>	0.88	<b>0.92</b>	0.93	<b>1.88</b>	1.95
	Max	<b>1.12</b>	1.14	1.20	<b>1.15</b>	<b>2.82</b>	3.49
	StD	0.13	0.14	0.15	0.14	0.60	0.79
16	Min	0.85	<b>0.75</b>	0.82	<b>0.77</b>	1.17	<b>1.06</b>
	Avg	<b>0.98</b>	1.00	<b>1.04</b>	1.05	<b>1.97</b>	2.08
	Max	<b>1.16</b>	1.31	1.31	<b>1.28</b>	<b>3.14</b>	3.69
	StD	0.09	0.17	0.12	0.13	0.65	0.83
24	Min	<b>0.91</b>	0.93	1.02	<b>0.98</b>	1.16	<b>1.09</b>
	Avg	<b>1.15</b>	1.16	1.22	<b>1.19</b>	<b>2.06</b>	2.19
	Max	1.72	<b>1.60</b>	1.81	<b>1.61</b>	<b>3.49</b>	4.08
	StD	0.20	0.21	0.18	0.16	0.70	0.91
32	Min	1.05	<b>1.04</b>	<b>1.07</b>	1.12	1.33	<b>1.26</b>
	Avg	1.51	<b>1.49</b>	1.54	<b>1.51</b>	<b>2.24</b>	2.41
	Max	<b>2.52</b>	2.63	<b>2.52</b>	2.62	<b>3.71</b>	4.51
	StD	0.45	0.45	0.42	0.43	0.74	1.02

the characteristics of the benchmarks). For the average and maximum overhead ratio, local scheduling is always better than batched scheduling. The reader can observe on Appendix. ?? that the maximum overhead values for 8, 16, 24 and 32 threads are given by the pyramid benchmark in the path right set.

In summary, we can say that both the local and batched scheduling strategies have similar overhead results on worst case scenarios for the  $NS$ ,  $SS_{LF_2}$  and  $FS_{LF_2+PCC}$  designs.

## 5 Conclusions and Further Work

### Acknowledgments

### References

1. Areias, M., Rocha, R.: An Efficient and Scalable Memory Allocator for Multi-threaded Tabled Evaluation of Logic Programs. In: International Conference on Parallel and Distributed Systems. pp. 636–643. IEEE Computer Society (2012)

2. Areias, M., Rocha, R.: A lock-free hash trie design for concurrent tabled logic programs. *International Journal of Parallel Programming* pp. 1–21 (2015)
3. Freire, J., Swift, T., Warren, D.S.: Beyond Depth-First: Improving Tabled Logic Programs through Alternative Scheduling Strategies. In: *International Symposium on Programming Language Implementation and Logic Programming*. pp. 243–258. No. 1140 in LNCS, Springer (1996)
4. Sagonas, K., Swift, T.: An Abstract Machine for Tabled Execution of Fixed-Order Stratified Logic Programs. *ACM Transactions on Programming Languages and Systems* 20(3), 586–634 (1998)