

On Extending a Full-Sharing Multithreaded Tabling Design with Batched Scheduling

Miguel Areias and Ricardo Rocha

CRACS & INESC-TEC LA

Faculty of Sciences, University of Porto, Portugal

miguel-areias@dcc.fc.up.pt ricroc@dcc.fc.up.pt

Yap Prolog: *<http://www.dcc.fc.up.pt/~vsc/Yap>*

Project SIBILA: *<http://cracs.fc.up.pt/>*



Prolog and SLD Resolution

- Prolog systems are known to have good performances and flexibility, but they are based on SLD resolution, which limits the potential of the Logic Programming paradigm.
- SLD resolution cannot deal properly with the following situations:
 - ◆ **Positive Infinite Cycles** (insufficient expressiveness)
 - ◆ **Negative Infinite Cycles** (inconsistence)
 - ◆ **Redundant Computations** (inefficiency)

SLD Resolution: Infinite Cycles

```
c1)    a(X) :- b(X).  
c2)    a(2).  
  
c3)    b(X) :- a(X).  
c4)    b(1).
```

1.a(X)

SLD Resolution: Infinite Cycles

```
c1)    a(X) :- b(X).  
c2)    a(2).  
  
c3)    b(X) :- a(X).  
c4)    b(1).
```

```
1.a(X)  
  | c1  
2.b(X)
```

SLD Resolution: Infinite Cycles

```
c1)    a(X) :- b(X).  
c2)    a(2).  
  
c3)    b(X) :- a(X).  
c4)    b(1).
```

```
1.a(X)  
  | c1  
2.b(X)  
  | c3  
3.a(X)
```

SLD Resolution: Infinite Cycles

```
c1)    a(X) :- b(X).  
c2)    a(2).  
  
c3)    b(X) :- a(X).  
c4)    b(1).
```

```
1.a(X)  
  | c1  
2.b(X)  
  | c3  
3.a(X)  
  | c1
```

Infinite Cycle

SLD Resolution: Infinite Cycles

c1) a(X) :- b(X).

c2) a(2).

c3) b(X) :- a(X).

c4) b(1).

1.a(X)

c1

2.b(X)

c3

3.a(X)

c1

Infinite Cycle

Tabling in Prolog Systems

- **Tabling** is an **implementation technique** that **overcomes** some of the **limitations** of **Prolog** systems:
 - ◆ Tabled subgoals are evaluated by storing their answers in an appropriate data space, called the **table space**.
 - ◆ Repeated calls to tabled subgoals are resolved by **consuming** the answers already stored in the table instead of **being re-evaluated** against the program clauses.

Tabling in Prolog Systems

- **Tabling** is an **implementation technique** that **overcomes** some of the **limitations** of **Prolog** systems:
 - ◆ Tabled subgoals are evaluated by storing their answers in an appropriate data space, called the **table space**.
 - ◆ Repeated calls to tabled subgoals are resolved by **consuming** the answers already stored in the table instead of **being re-evaluated** against the program clauses.
- Implementations of **Tabling** are currently available in systems like:
 - ◆ XSB Prolog, **Yap Prolog**, B-Prolog, ALS-Prolog, Mercury, Ciao Prolog and more recently Picat.
- **Multithreading** combined with **Tabling**:
 - ◆ XSB Prolog
 - ◆ **YapTab-Mt** [ICLP 2012].

YapTab-Mt - Advantages

- An **abstraction layer** with **high-level constructors** that provide access to the **dynamic programming (tabling)** support:
 - ◆ Instruction: **`:- table predicate/arity.`**
 - ◆ Scheduling: **`:- tabling_mode(predicate, batched).`**

YapTab-Mt - Advantages

- An **abstraction layer** with **high-level constructors** that provide access to the **dynamic programming (tabling)** support:
 - ◆ Instruction: **`:- table predicate/arity.`**
 - ◆ Scheduling: **`:- tabling_mode(predicate, batched).`**
- **Thread API** is **POSIX Threads compliant**:
 - ◆ **Management** - creating, joining , yielding, etc.
 - ◆ **Monitoring** - statistics, properties, etc.
 - ◆ **Synchronization** - mutex creation, statistics, etc.

YapTab-Mt - Advantages

- An **abstraction layer** with **high-level constructors** that provide access to the **dynamic programming (tabling)** support:
 - ◆ Instruction: **`:- table predicate/arity.`**
 - ◆ Scheduling: **`:- tabling_mode(predicate, batched).`**
- **Thread API** is **POSIX Threads compliant**:
 - ◆ **Management** - creating, joining , yielding, etc.
 - ◆ **Monitoring** - statistics, properties, etc.
 - ◆ **Synchronization** - mutex creation, statistics, etc.
- Write complex **dynamic programming** applications using the **Prolog** programming language.
 - ◆ **Procedures** in **Prolog** can be written as **logical specifications**, which are closer to **mathematical notation**.

Table Space - Internal Architecture

➤ **Table Entry**: stores generic about the predicates.

◆ **:-table predicate/2.**

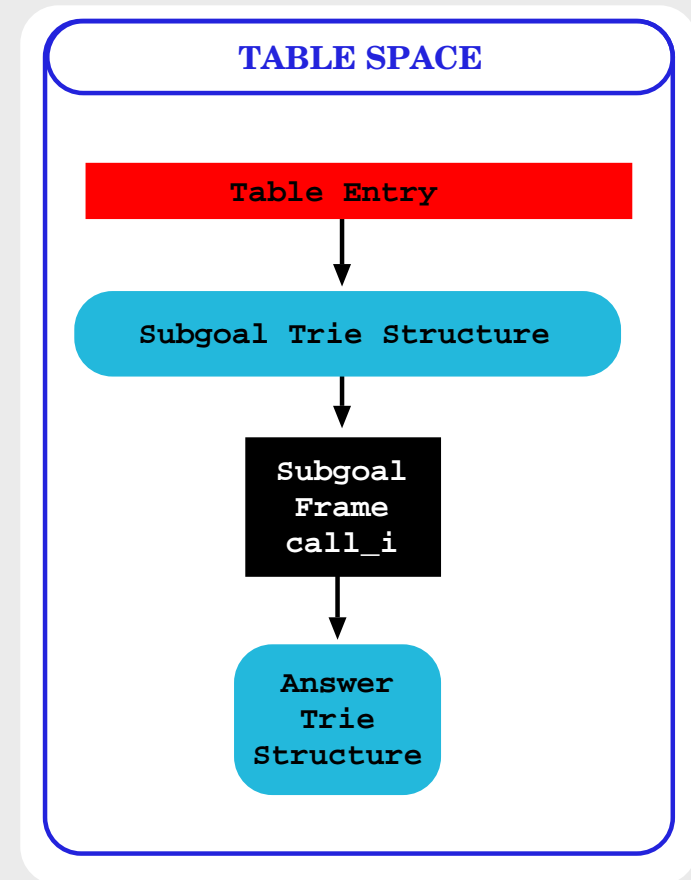


Table Space - Internal Architecture

- **Table Entry**: stores generic about the predicates.
 - ◆ **`:-table predicate/2`**.
- **Subgoal Trie Structure**: stores the **identifier** of the computations.
 - ◆ **`predicate(computation_id, Answer)`**.




Table Space - Internal Architecture

- **Table Entry**: stores generic about the predicates.
 - ◆ `:-table predicate/2`.
- **Subgoal Trie Structure**: stores the **identifier** of the computations.
 - ◆ `predicate(computation_id, Answer)`.
- **Answer Trie Structure**: stores the **answers** of the computations.
 - ◆ `predicate(computation_id, Answer)`.

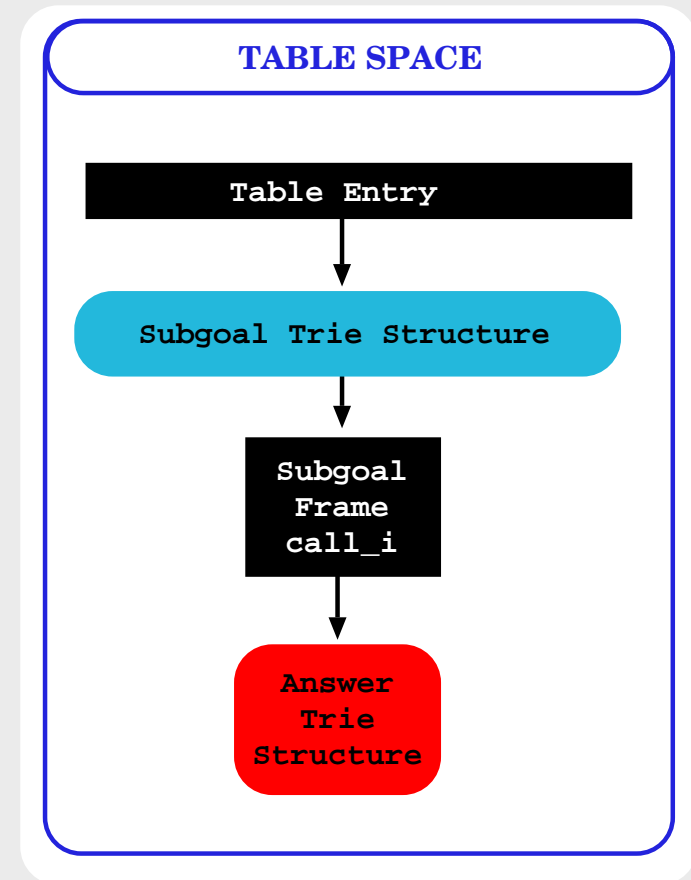


Table Space - Example

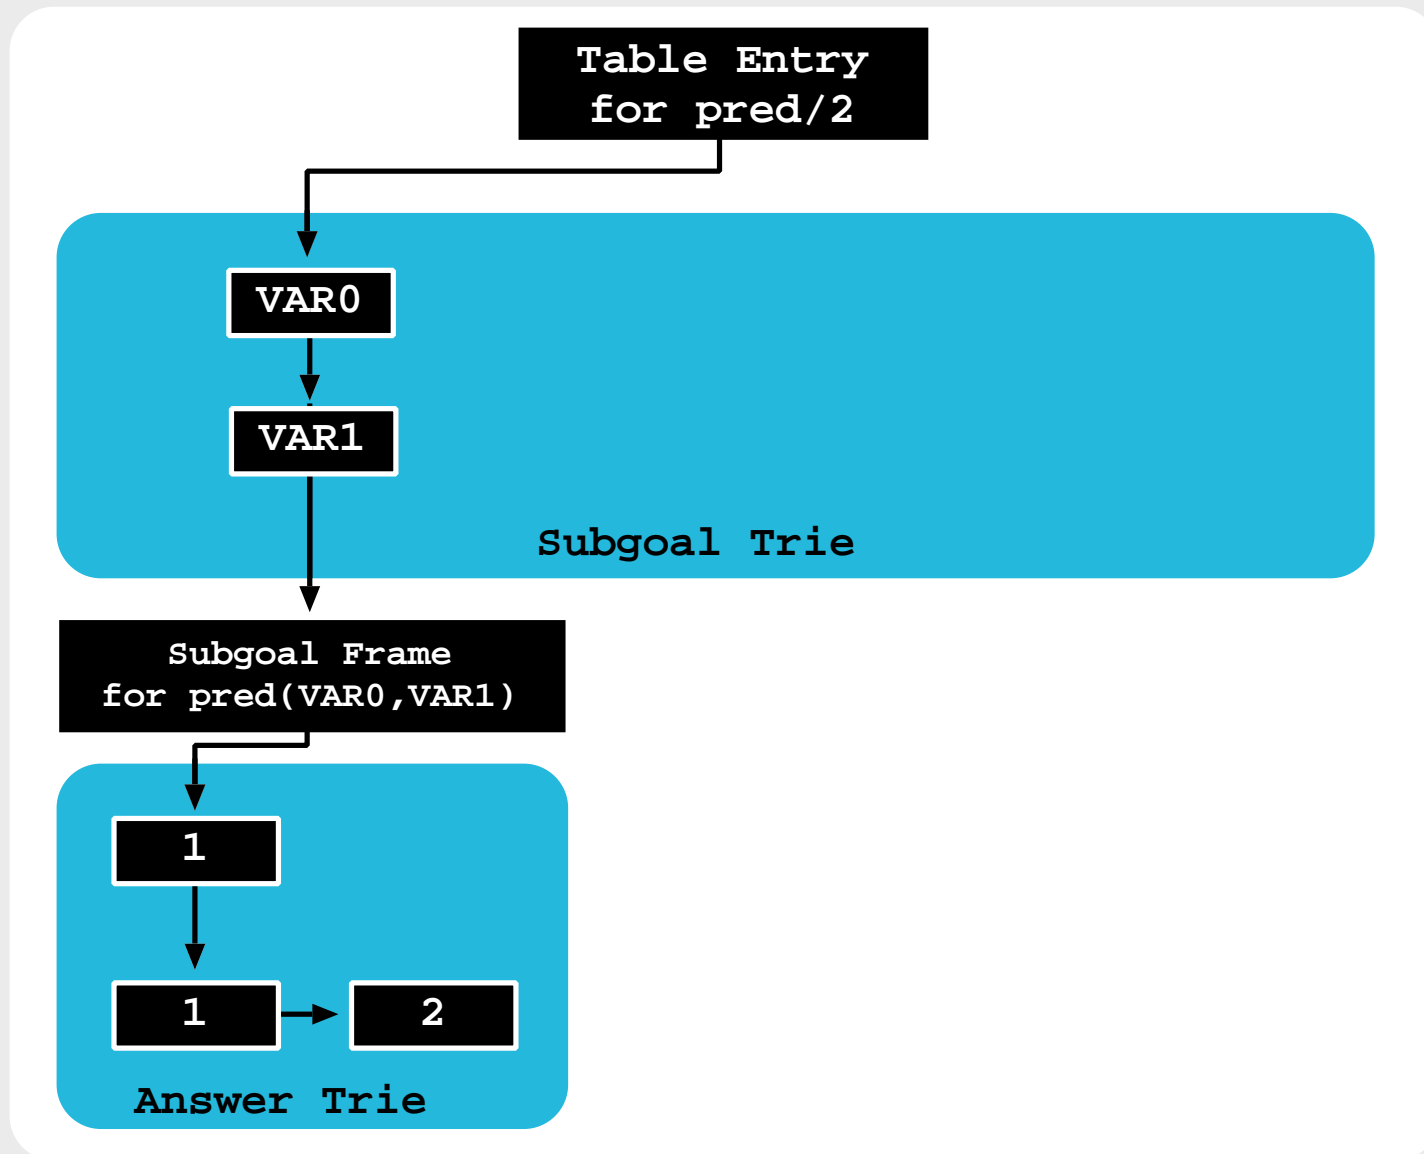
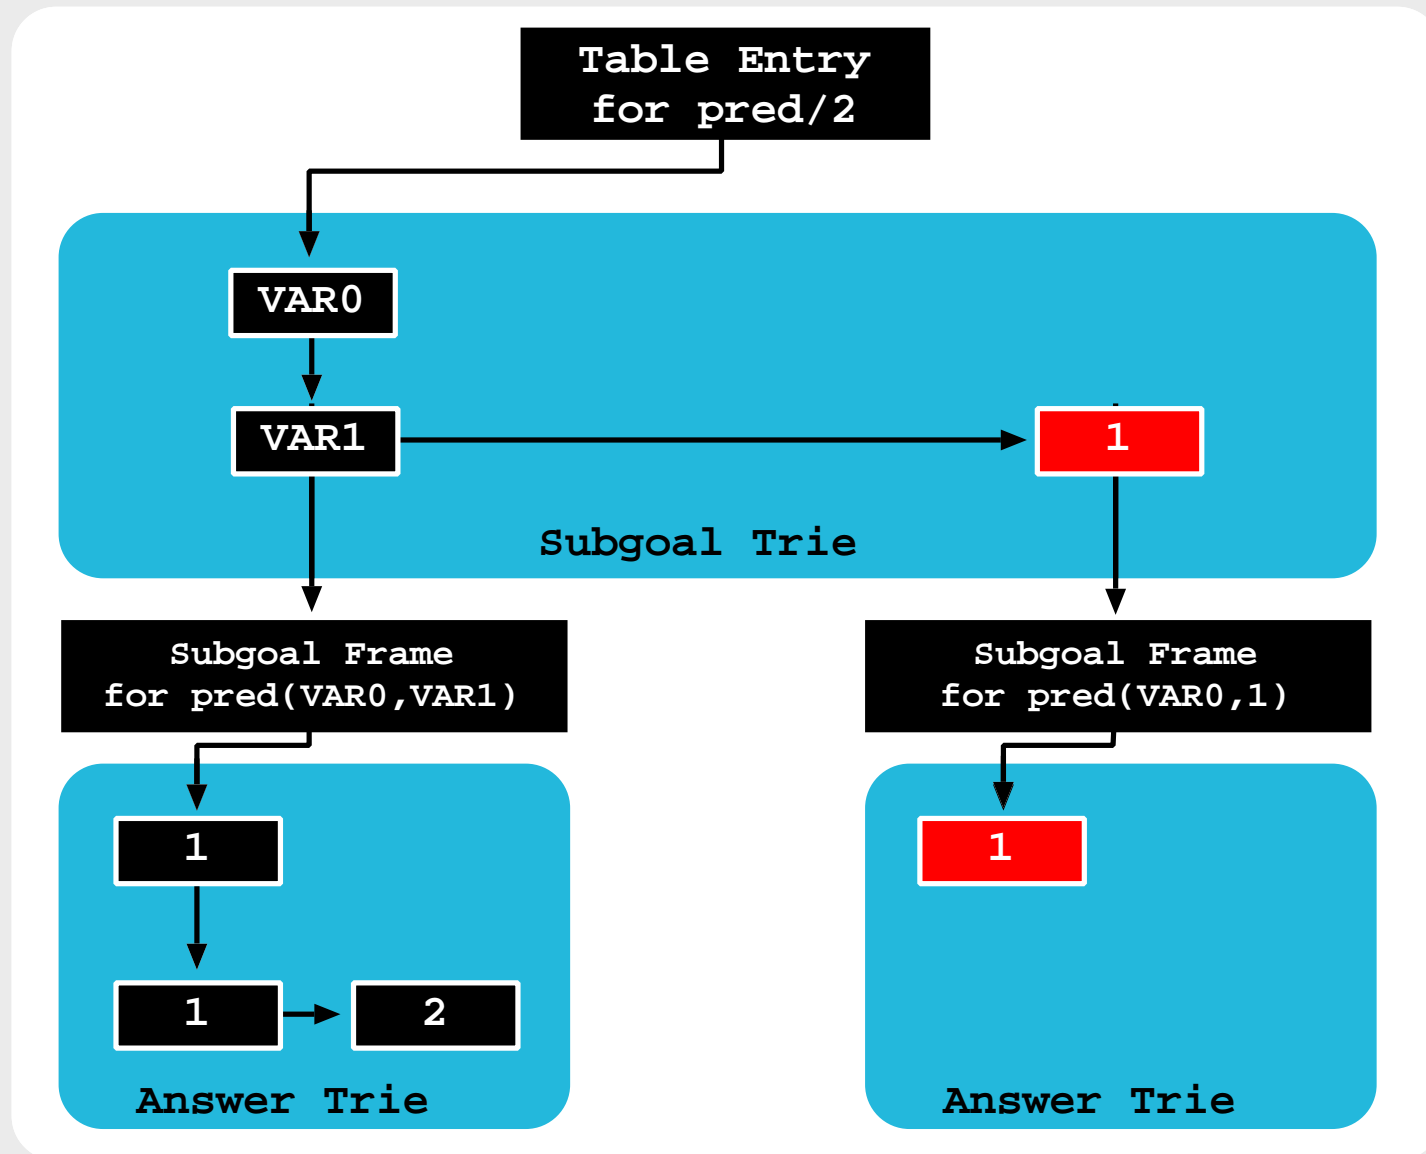


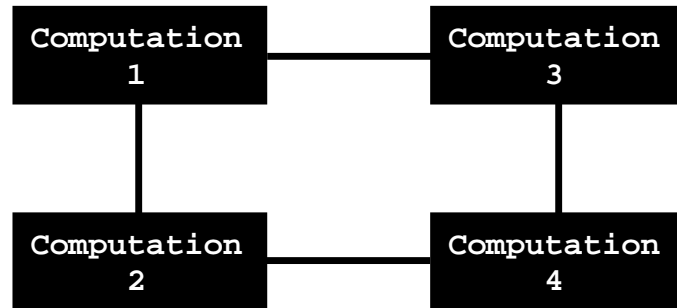
Table Space - Example



Tabling Scheduling Strategies - Local vs Batched

LOCAL SCHEDULING

Cluster of Dependent Computations

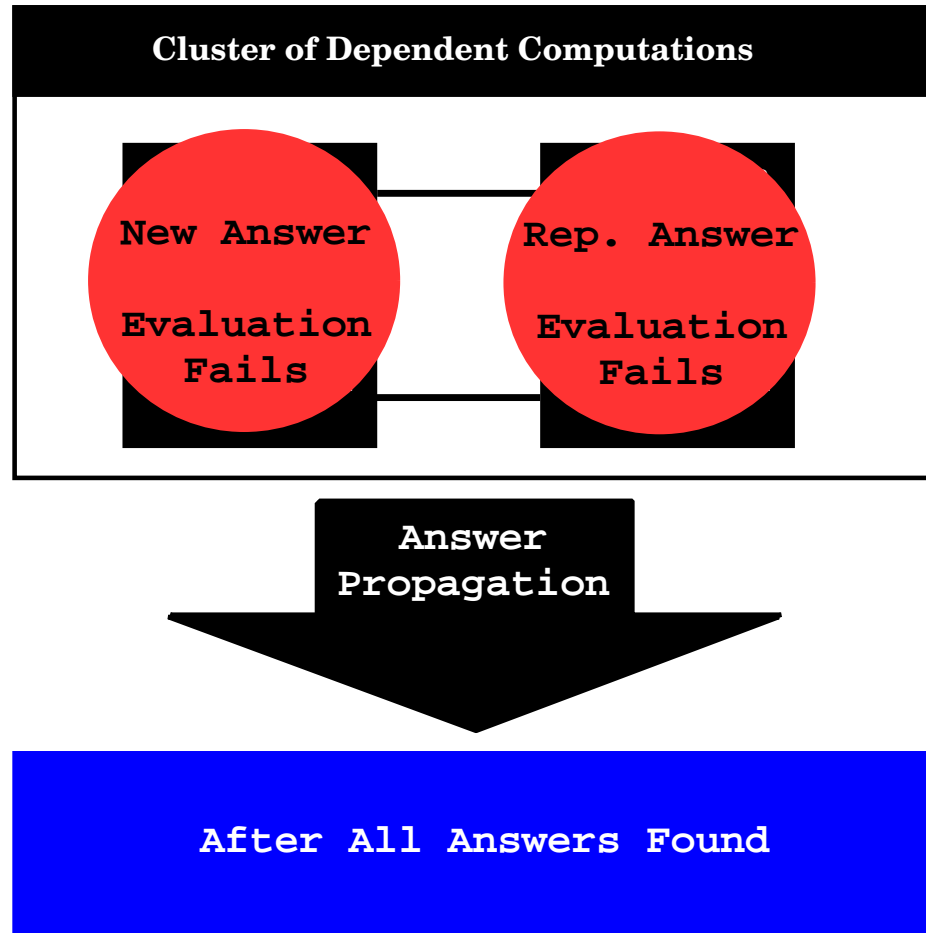


Answer
Propagation

After All Answers Found

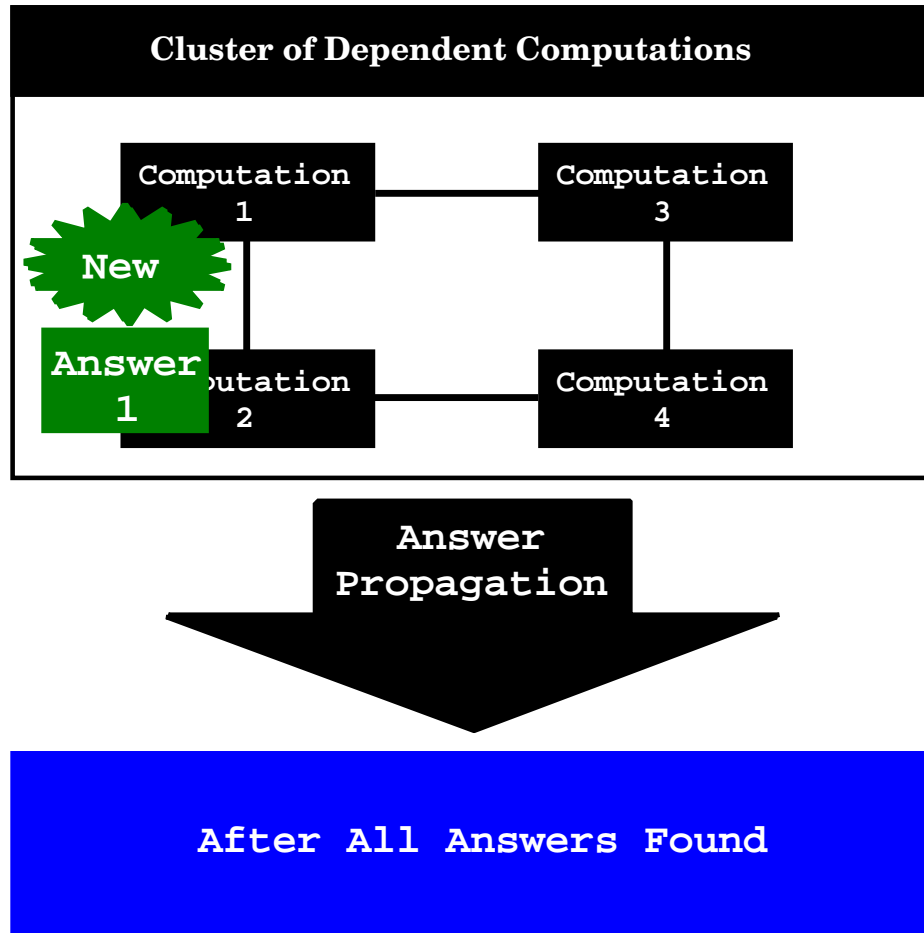
Tabling Scheduling Strategies - Local vs Batched

LOCAL SCHEDULING



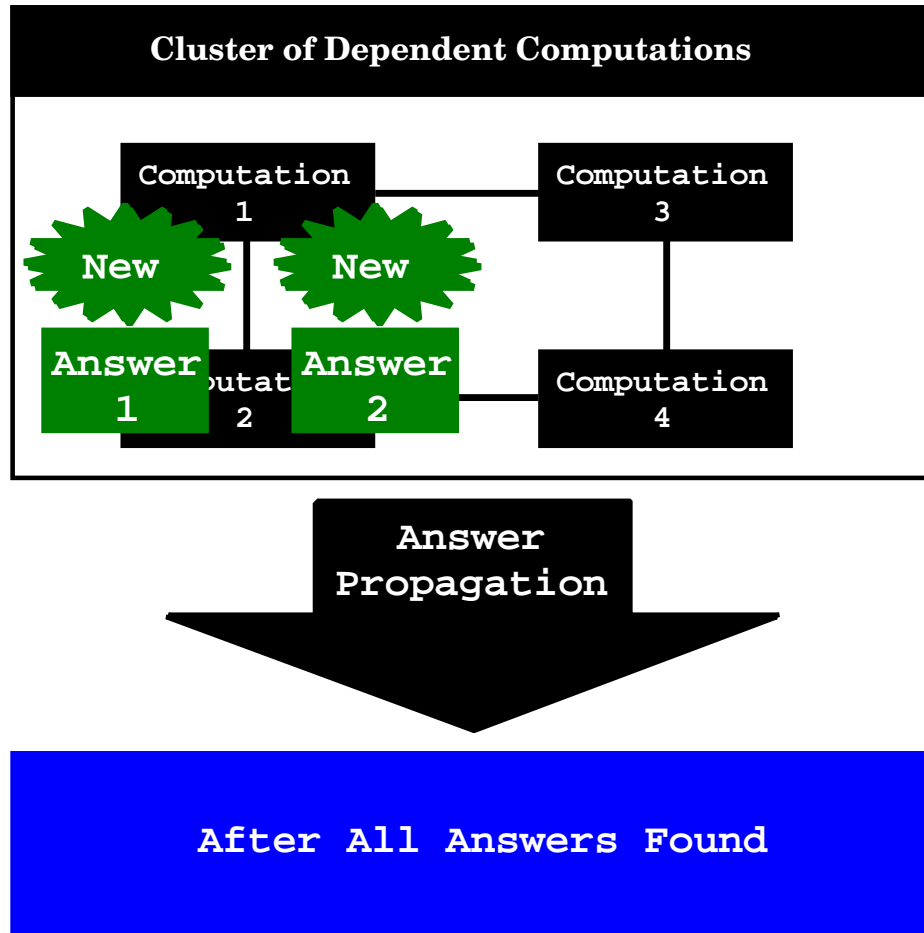
Tabling Scheduling Strategies - Local vs Batched

LOCAL SCHEDULING



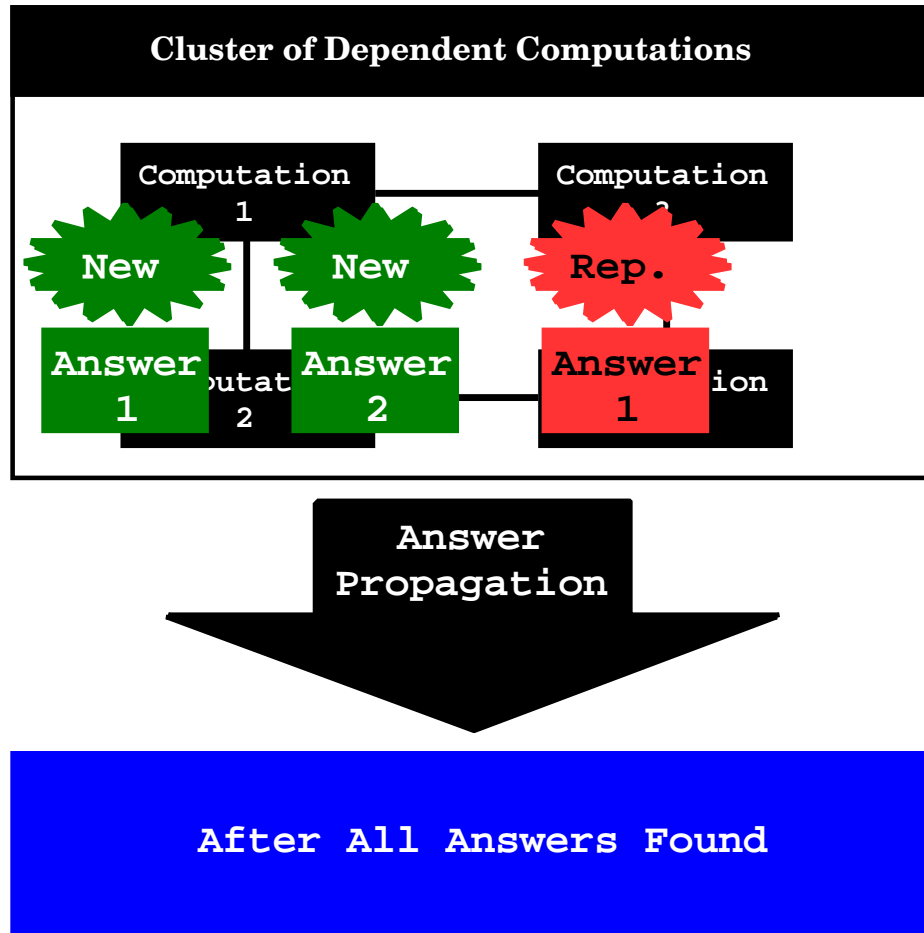
Tabling Scheduling Strategies - Local vs Batched

LOCAL SCHEDULING



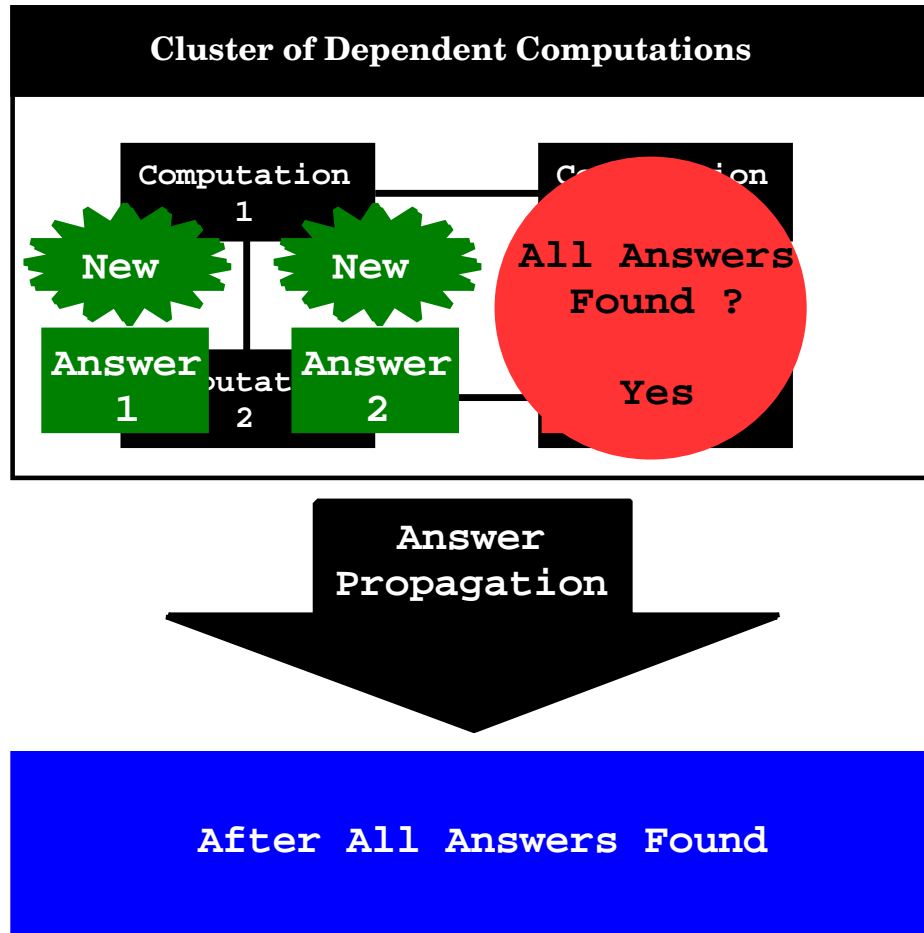
Tabling Scheduling Strategies - Local vs Batched

LOCAL SCHEDULING



Tabling Scheduling Strategies - Local vs Batched

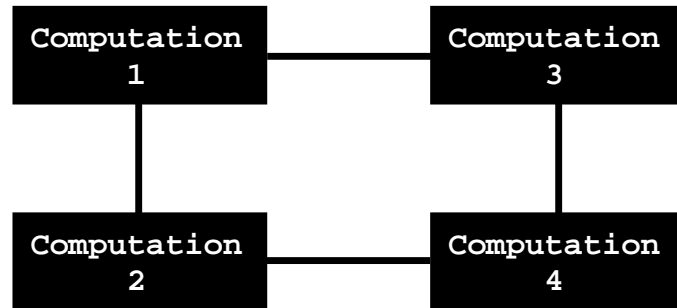
LOCAL SCHEDULING



Tabling Scheduling Strategies - Local vs Batched

LOCAL SCHEDULING

Cluster of Dependent Computations




Answer
Propagation

Answer 1 After Answer 2 Answers Found

Tabling Scheduling Strategies - Local vs Batched

LOCAL SCHEDULING

Cluster of Dependent Computations




Answer
Propagation

After All Answers Found

BATCHED SCHEDULING

Cluster of Dependent Computations



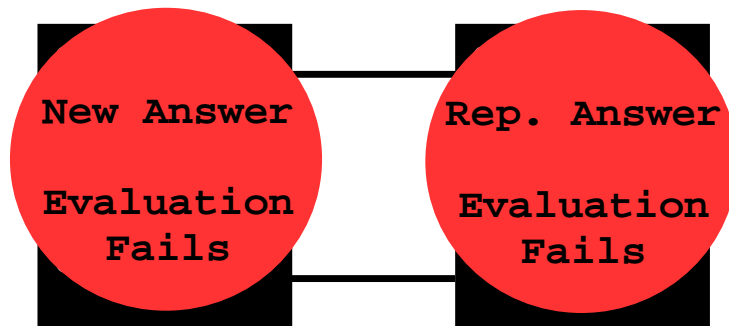
Answer
Propagation

Whenever a New Answer is Found

Tabling Scheduling Strategies - Local vs Batched

LOCAL SCHEDULING

Cluster of Dependent Computations

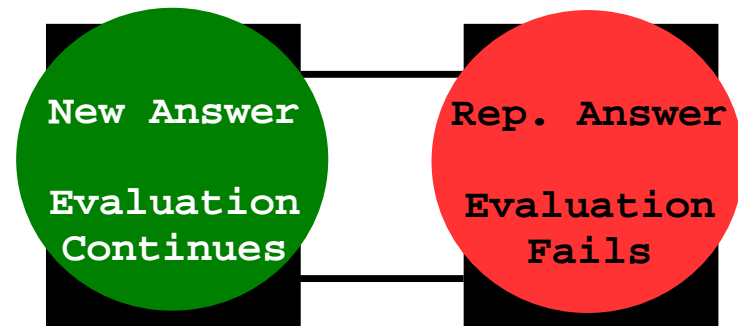


Answer
Propagation

After All Answers Found

BATCHED SCHEDULING

Cluster of Dependent Computations



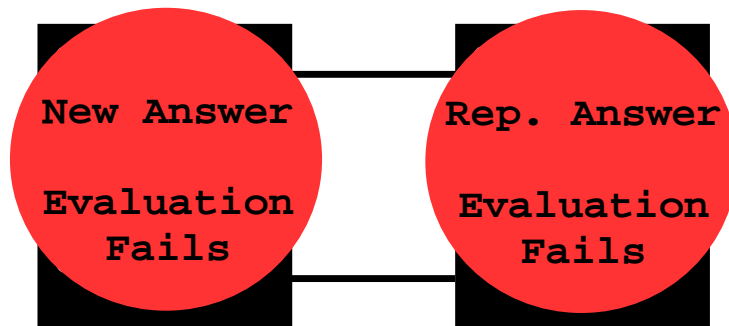
Answer
Propagation

Whenever a New Answer is Found

Tabling Scheduling Strategies - Local vs Batched

LOCAL SCHEDULING

Cluster of Dependent Computations




Answer
Propagation

After All Answers Found

BATCHED SCHEDULING

Cluster of Dependent Computations




Answer
Propagation

Whenever a New Answer is Found

Tabling Scheduling Strategies - Local vs Batched

LOCAL SCHEDULING

Cluster of Dependent Computations




Answer
Propagation

After All Answers Found

BATCHED SCHEDULING

Cluster of Dependent Computations



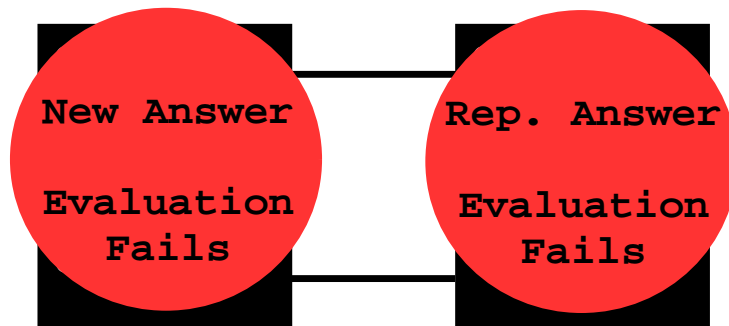
Answer
Propagation

Answer 1 never a New Answer is Found

Tabling Scheduling Strategies - Local vs Batched

LOCAL SCHEDULING

Cluster of Dependent Computations




Answer
Propagation

After All Answers Found

BATCHED SCHEDULING

Cluster of Dependent Computations



Answer
Propagation

Answer 1 never Answer 2 Answer is Found

Tabling Scheduling Strategies - Local vs Batched

LOCAL SCHEDULING

Cluster of Dependent Computations

New Answer
Evaluation
Fails

Rep. Answer
Evaluation
Fails

Answer
Propagation

After All Answers Found

BATCHED SCHEDULING

Cluster of Dependent Computations

Computation
1

Computation
2

Rep.

Computation
2

Answer
1


Answer
Propagation

Answer
1


Answer
2

Answer is Found


YapTab-Mt - Internal Architecture



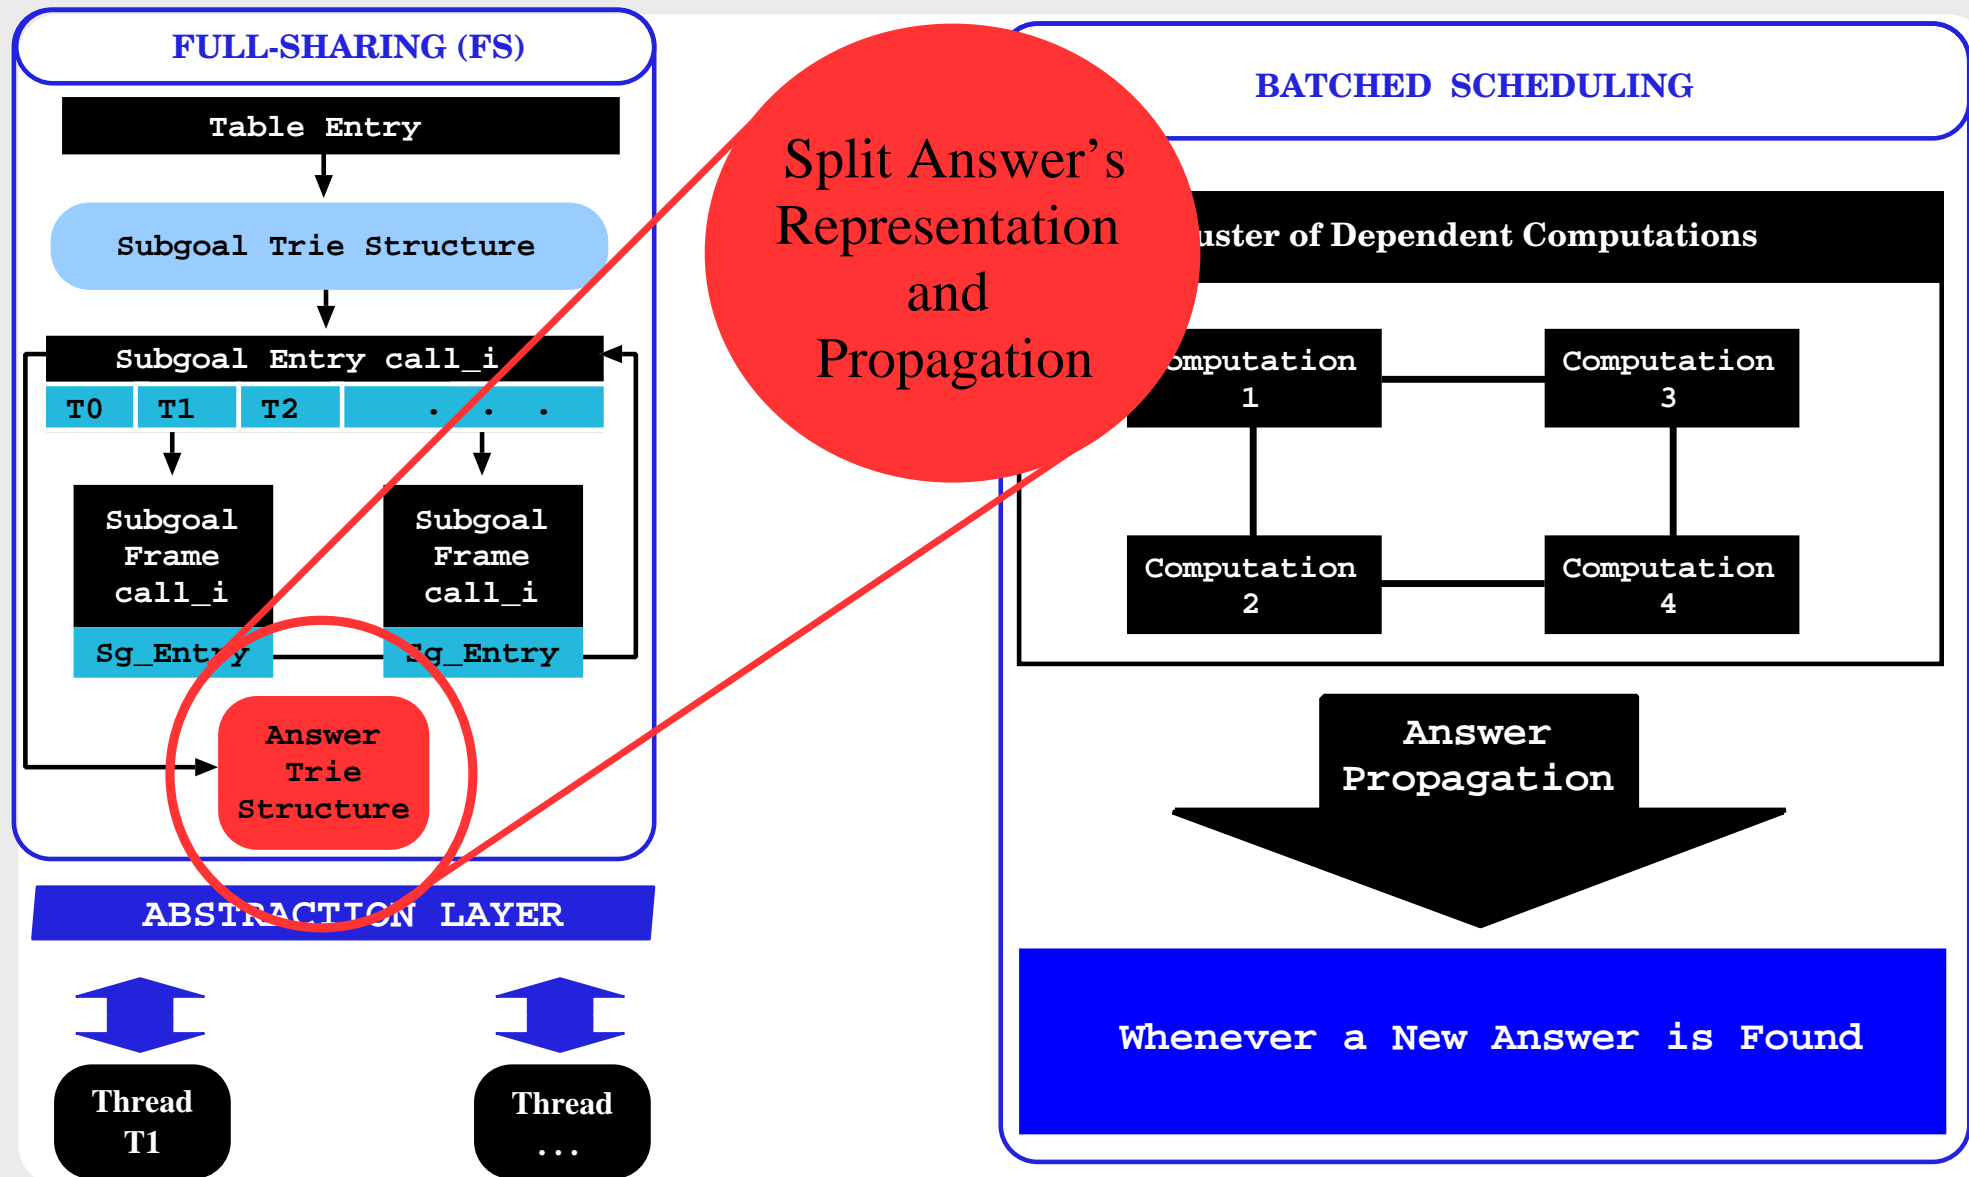
YapTab-Mt - Internal Architecture



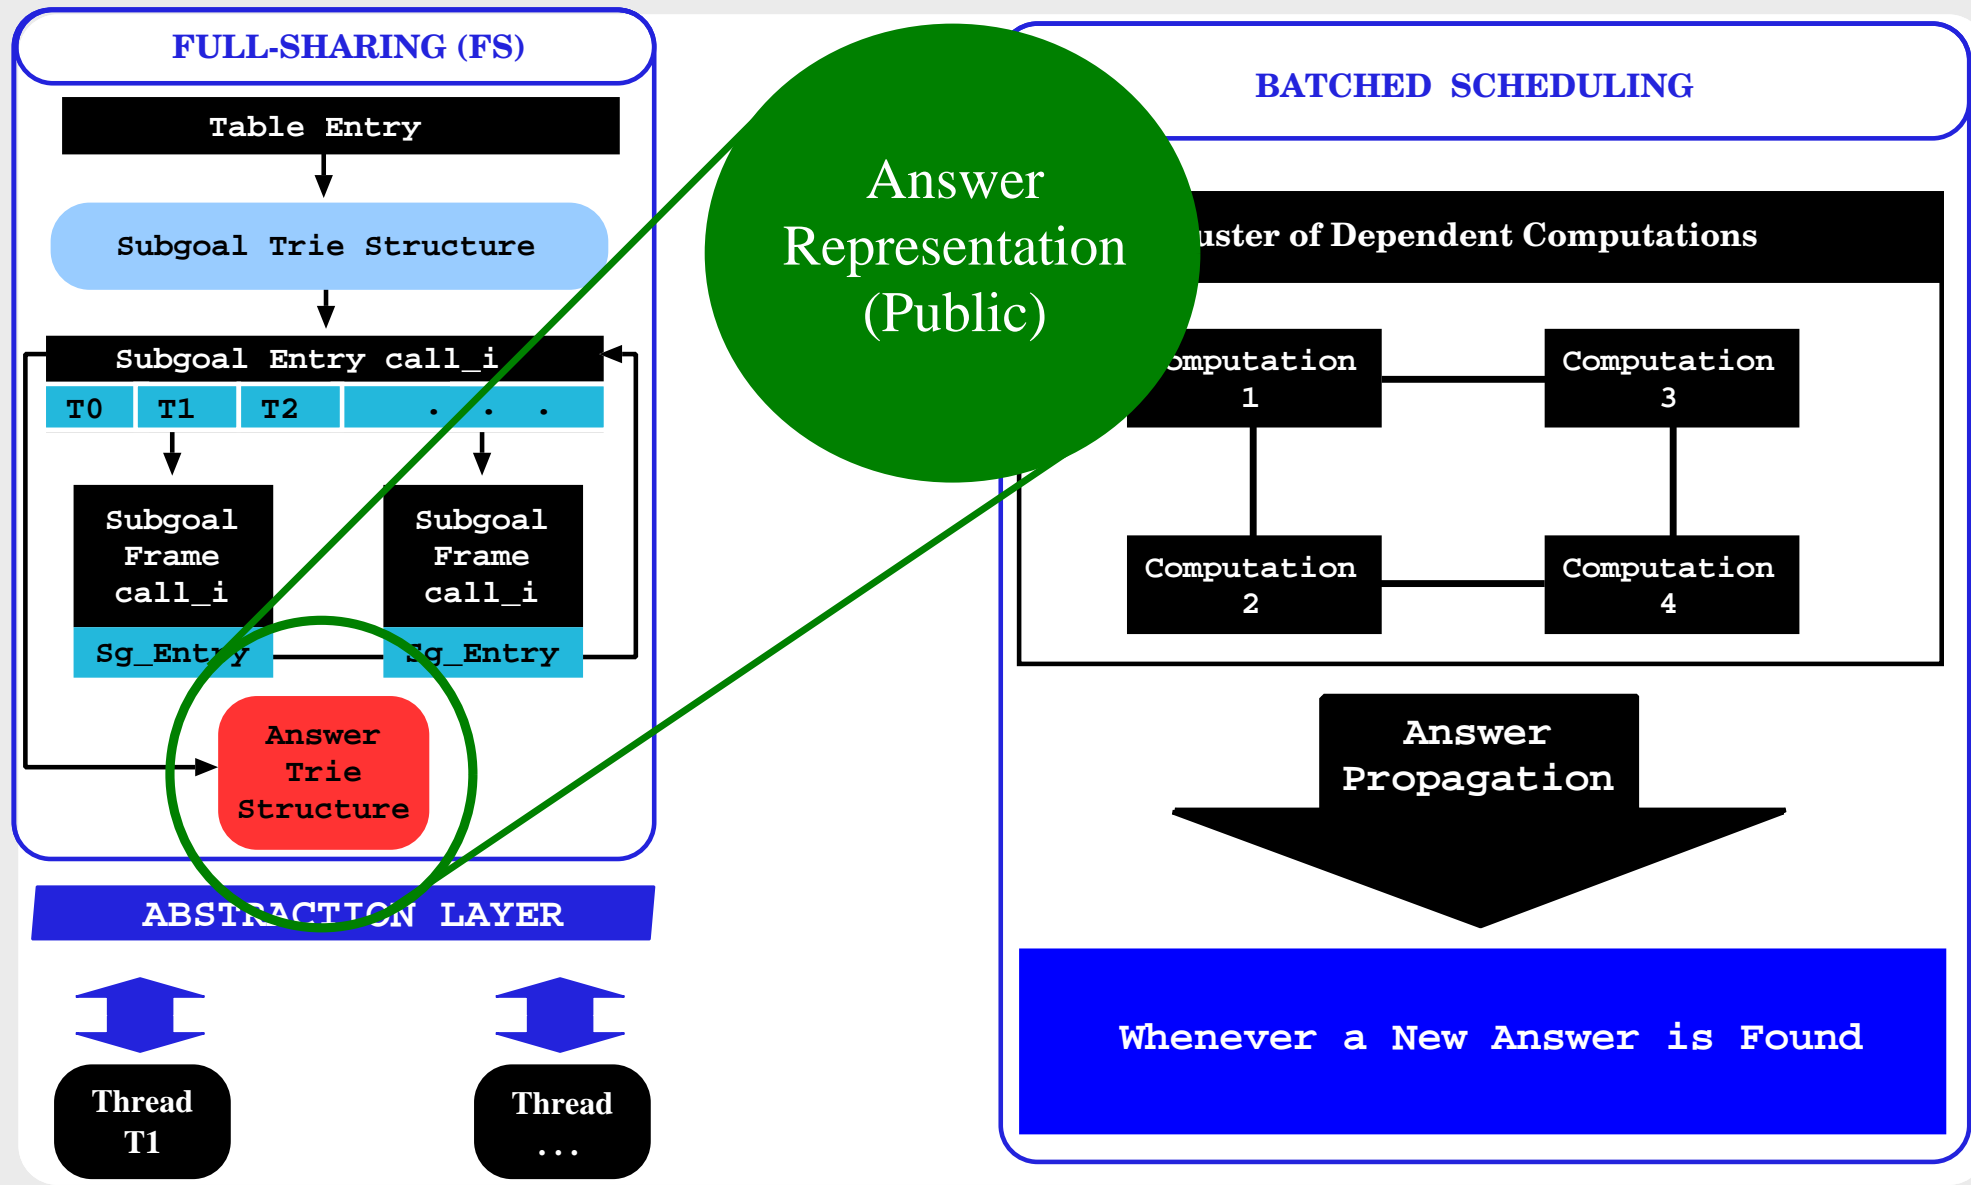
YapTab-Mt - Internal Architecture



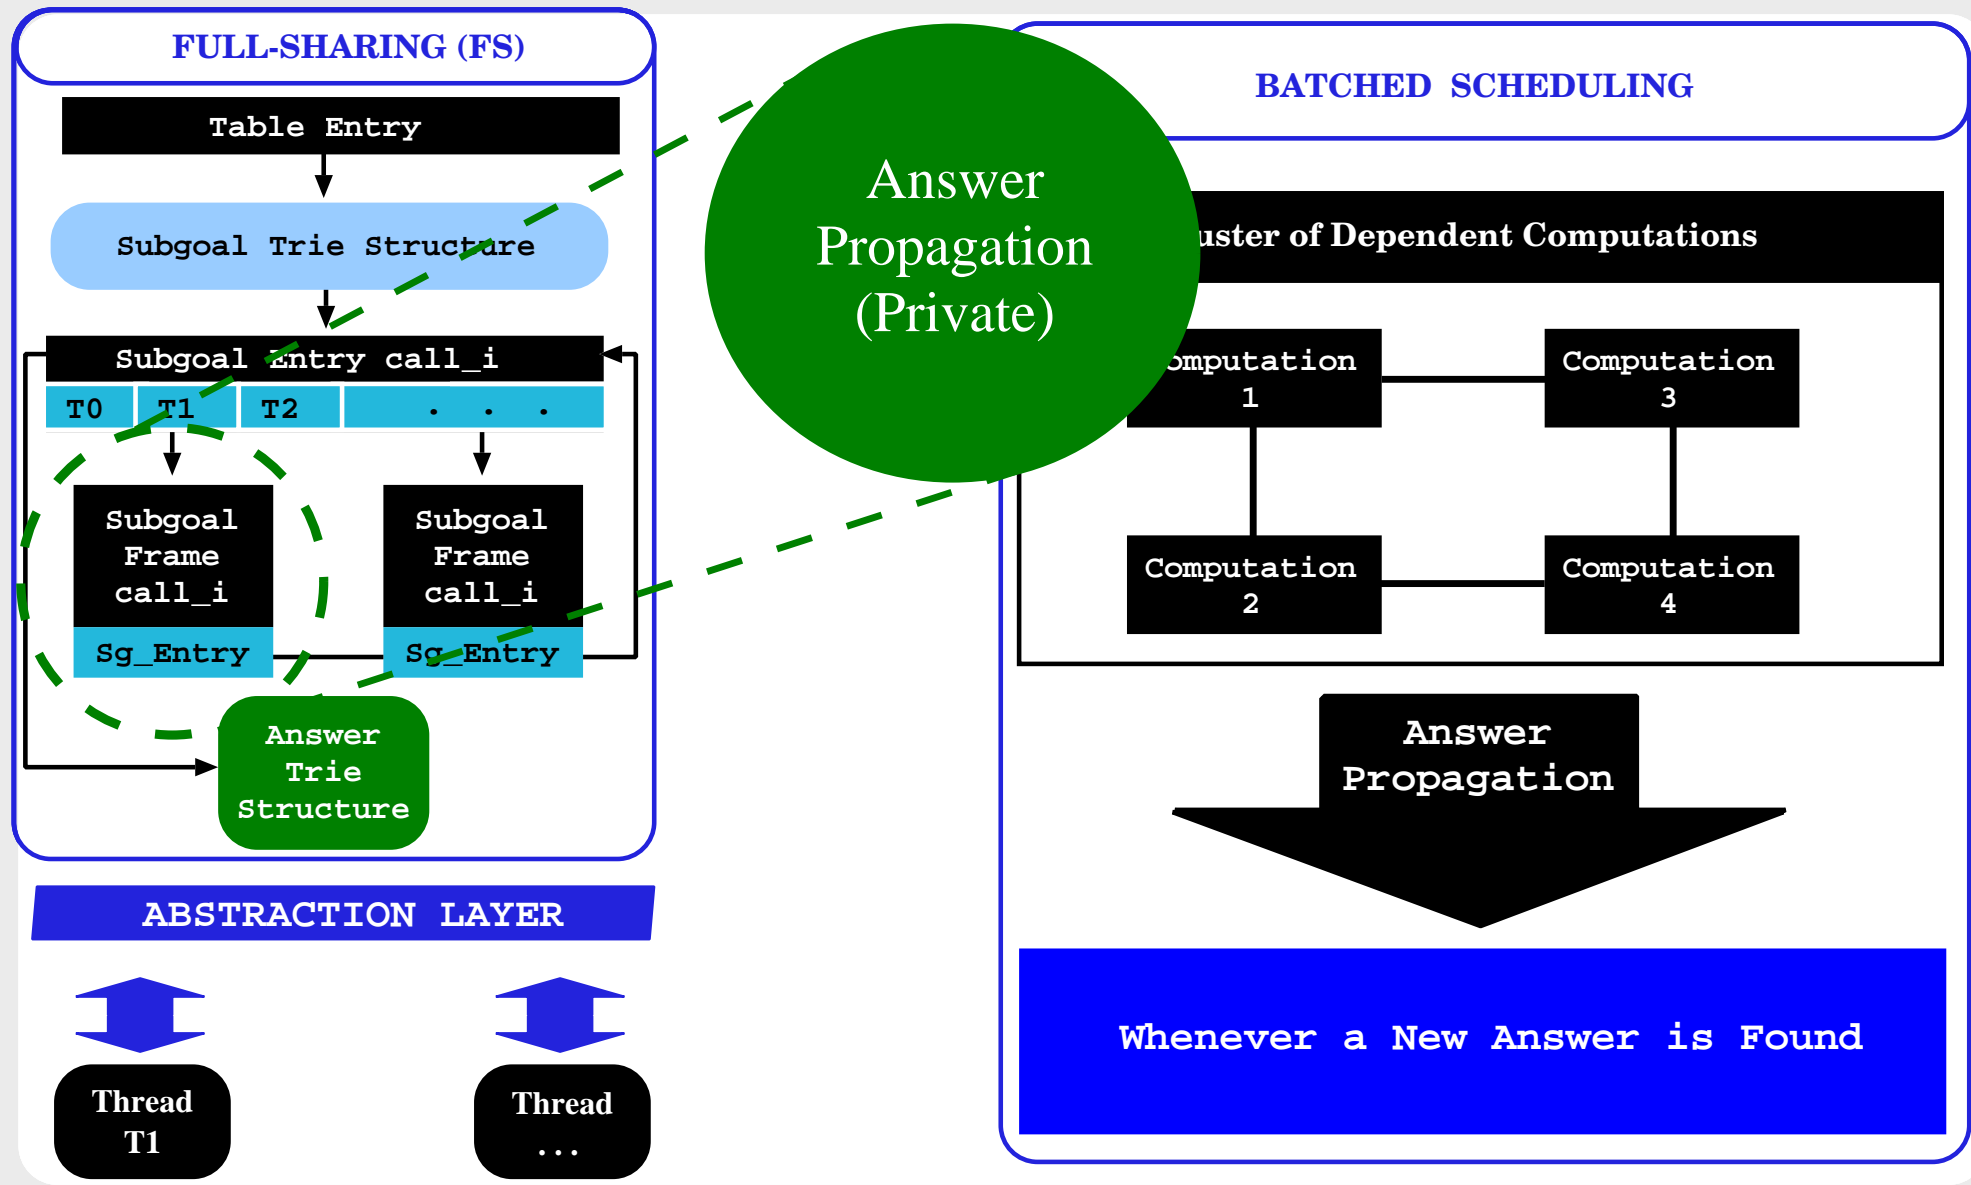
Private Answer Chaining - Key Idea



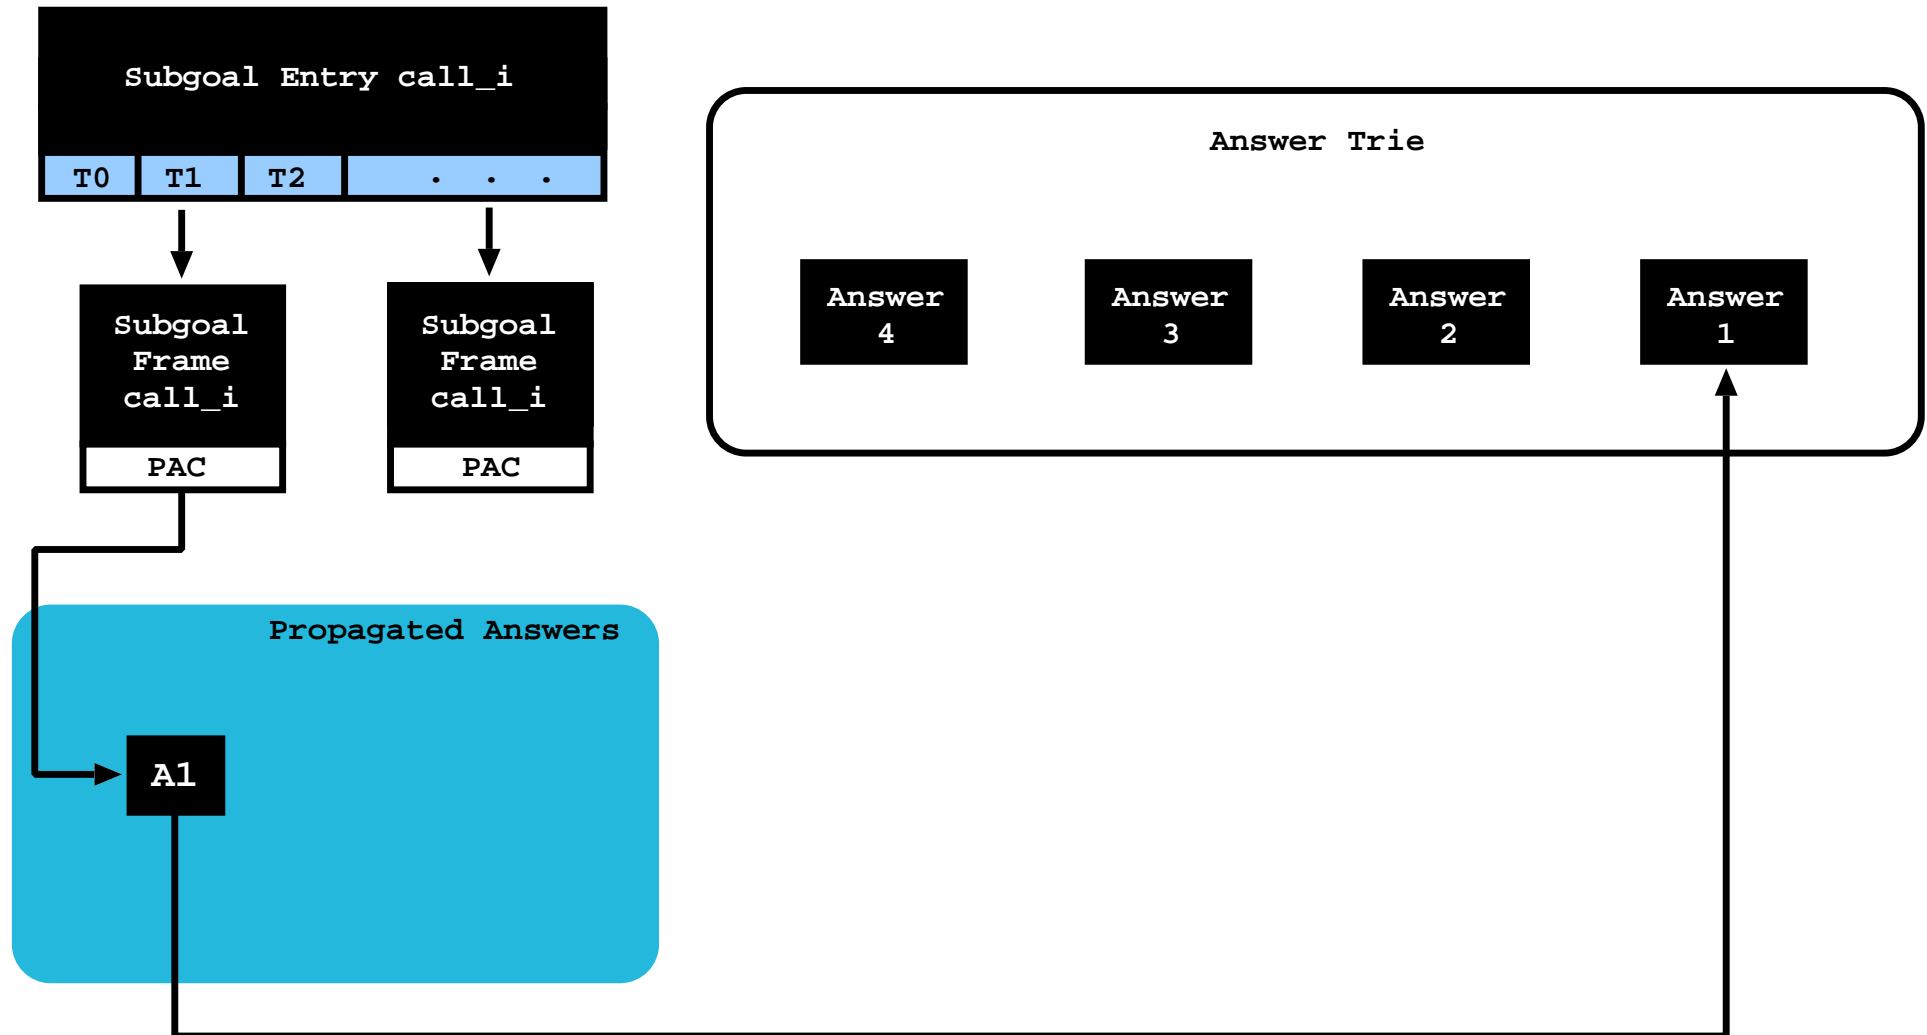
Private Answer Chaining - Key Idea




Private Answer Chaining - Key Idea



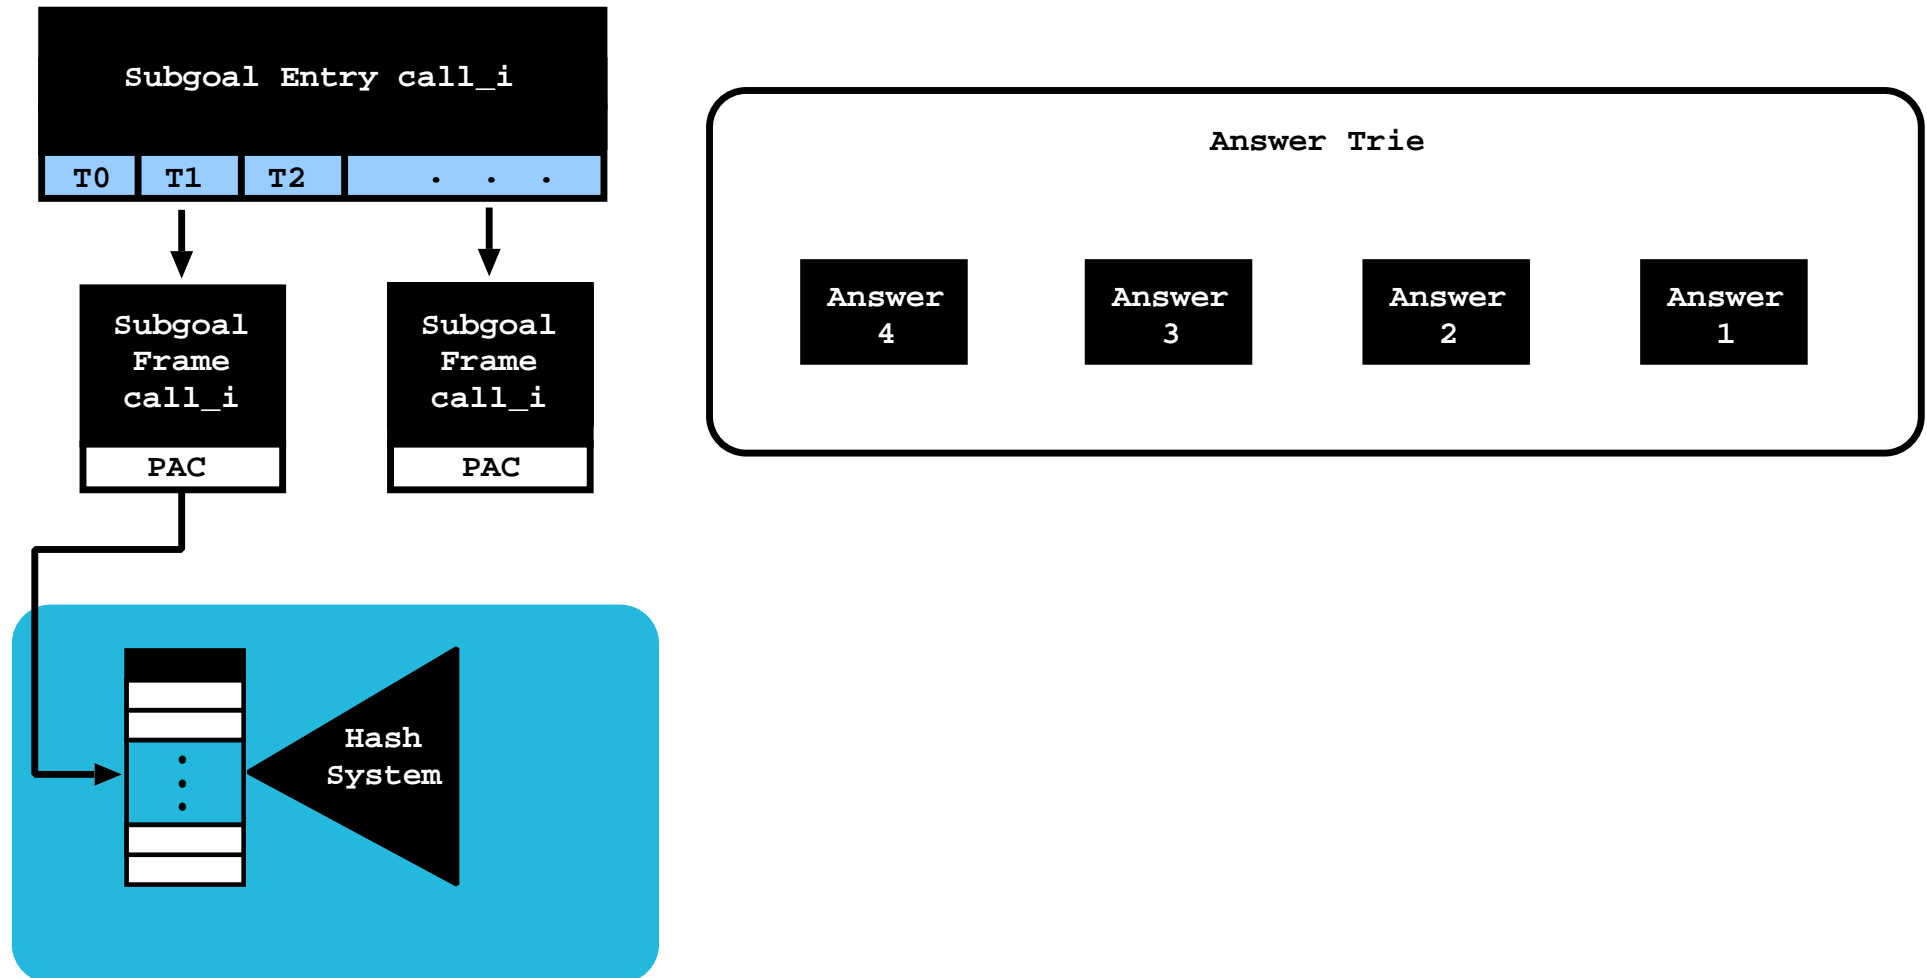
Private Answer Chaining - Internals



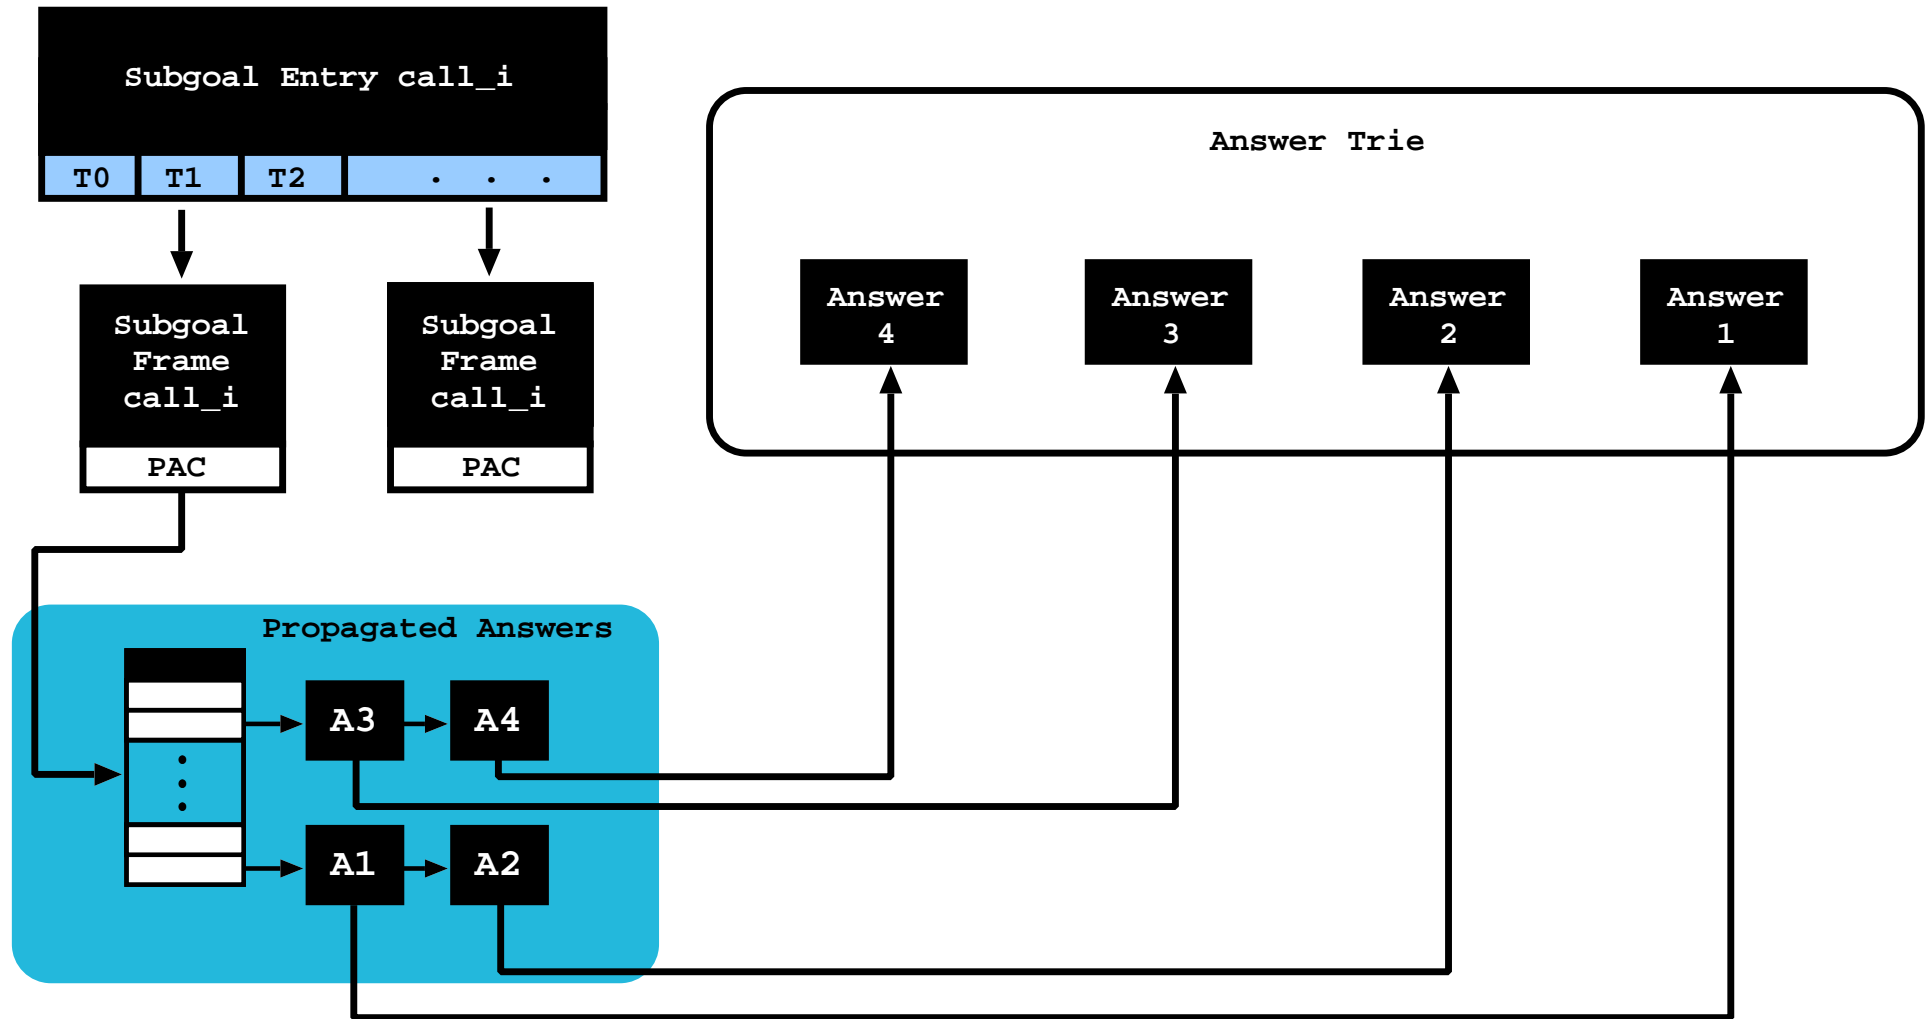
Private Answer Chaining - Internals



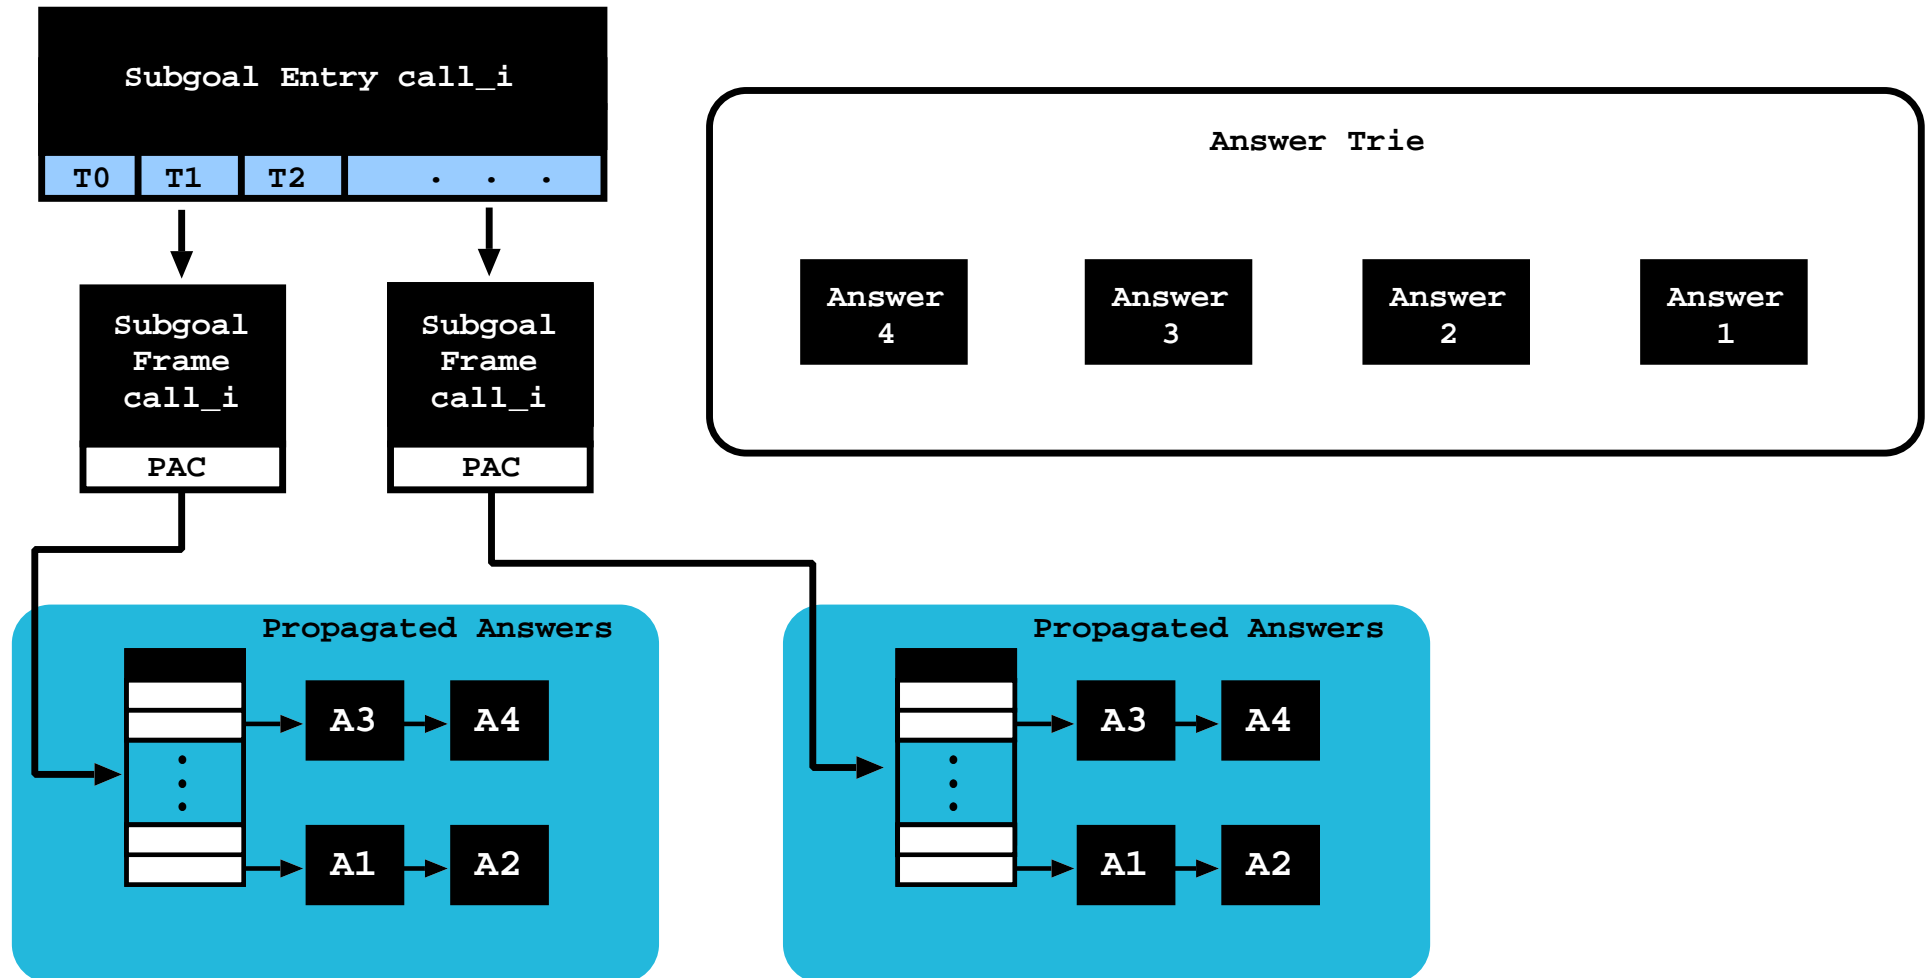
Private Answer Chaining - Internals



Private Answer Chaining - Internals



Private Answer Chaining - Internals



Experimental Results - Worst Case Scenarios

Threads		NS		FS	
		Local	Batched	Local	Batched
1	Min	0.53	0.55	1.01	0.95
	Avg	0.78	0.82	1.30	1.46
	Max	1.06	1.05	1.76	2.33
8	Min	0.66	0.63	1.16	0.99
	Avg	0.85	0.88	1.88	1.95
	Max	1.12	1.14	2.82	3.49
16	Min	0.85	0.75	1.17	1.06
	Avg	0.98	1.00	1.97	2.08
	Max	1.16	1.31	3.14	3.69
24	Min	0.91	0.93	1.16	1.09
	Avg	1.15	1.16	2.06	2.19
	Max	1.72	1.60	3.49	4.08
32	Min	1.05	1.04	1.33	1.26
	Avg	1.51	1.49	2.24	2.41
	Max	2.52	2.63	3.71	4.51

Conclusions and Further Work

- **Batched scheduling** can be an **useful strategy** in tabled logic programs that:
 - ◆ require an **eager propagation** of answers.
 - ◆ do not require the complete set of answers to be found (**partial evaluation** of the **search space**).


Conclusions and Further Work

- **Batched scheduling** can be an **useful strategy** in tabled logic programs that:
 - ◆ require an **eager propagation** of answers.
 - ◆ do not require the complete set of answers to be found (**partial evaluation** of the **search space**).
- We have presented the **PAC (Private Answer Chaining)** strategy that **enables** the **combination** of **batched scheduling** with **Full-Sharing**.
 - ◆ Split's **answers representation (public)** and **answer propagation (private)**.
- Experimental results showed that the **PAC** strategy is a **good first approach**, however there is still **room for improvement**.

Conclusions and Further Work

- **Batched scheduling** can be an **useful strategy** in tabled logic programs that:
 - ◆ require an **eager propagation** of answers.
 - ◆ do not require the complete set of answers to be found (**partial evaluation** of the **search space**).
- We have presented the **PAC (Private Answer Chaining)** strategy that **enables** the **combination** of **batched scheduling** with **Full-Sharing**.
 - ◆ Split's **answers representation (public)** and **answer propagation (private)**.
- Experimental results showed that the **PAC** strategy is a **good first approach**, however there is still **room for improvement**.
- Further work will include:
 - ◆ Use **timestamped tries** to **avoid** the search for the already propagated answers.
 - ◆ Seek for **new real-world problems** that allow us to improve and consolidate our framework.

Research Outline



Work related:



3 Journals



2 Book Series



3 Workshop Proceedings

Others:



2 Journals




0 Book Series






1 Workshop Proceeding




Research Outline



Work related:

-  3 Journals
-  2 Book Series
-  4 Workshop Proceedings

Others:

-  2 Journals
-  0 Book Series
-  1 Workshop Proceeding

Thank You !!!